Catalyst Works AI Governance Playbook 2026

Boubacar Barry, Catalyst Works Consulting Version 3.0 Final | April 16, 2026 | Classification: Internal, Founder Use Only Status: Phase 0 pre-revenue. Target: First $3,500 SHIELD Assessment invoiced by Day 30.12,000 to $16,500 total cash by Day 90.

This playbook integrates three source documents: the original Strategic Review, the AI Governance Practitioner Toolkit, and the prior CW AI Governance Playbook. It was pressure-tested through two Sankofa Council passes and a competitive intelligence review that captured the Fractional AI Officer category formation in March to April 2026. Every regulatory claim was verified against primary sources. Every pricing, positioning, and conversion mechanism was reviewed by five voices. The document is ordered for execution, not for reading.

Document Map

The document has three layers.

Layer 1: The Week 1 Action Pack. Twelve pages. Read every morning for thirty days. Contains only what is needed to close the first paid engagement. Everything else in the playbook is reference material.

Layer 2: The Strategy. The operating strategy for the practice. Positioning, offer stack, marketing and go-to-market, conversion mechanics, execution cadence, success criteria, deferred list. Consulted when a strategic question arises.

Layer 3: The Practitioner Toolkit. Internal working material. Discovery session scripts, diagnostic questions, maturity model, objection responses, business case, trigger events, content angles, failure modes, regulatory quick reference. Consulted during client engagements.

If a section of this document does not fit into one of these three layers, it is fluff and has been removed.

Layer 1: The Week 1 Action Pack

The only section of this document Boubacar reads daily for the first 30 days. Everything here exists to produce the first $3,500 SHIELD Assessment by Day 30 and 12,000 to $16,500 in total cash by Day 90.

1.1 The One Job

The single job of the next 30 days is a first paid SHIELD Assessment invoiced and collected. $3,500 in the account. Nothing else. Content gets published only if it directly supports that goal. Infrastructure gets built only if it directly supports that goal. Reading gets done only if it directly supports that goal.

The Sankofa Council's First Principles voice settled this. Every section of this Action Pack serves that one job. Every section outside this Action Pack is either reference material or future-state planning.

1.2 The 14-Day Task Sheet

Two actions per day for 14 days. No more, no less. Execute in order. If a day is missed, the next day does three tasks, not two.

Day 1 (Monday):

- Write the SHIELD one-pager (client-facing, two pages maximum, in the format specified in Section 2.5). Save to /mnt/user-data/outputs as a PDF ready to send.
- Build the list construction plan. Target: 30 qualified prospect names by end of Day 3. Use the method in Section 1.7.

Day 2 (Tuesday):

- Build the pre-session intake form (15 questions from the diagnostic, specified in Section 2.3). Deploy in Typeform or Google Forms. Link from the Calendly confirmation email.
- Update the catalystworks.consulting/governance landing page with the SHIELD Assessment at $3,500 publicly listed. Keep the 2,500 launch rate internal.

Day 3 (Wednesday):

- List construction executed. 30 names in the cold outreach spreadsheet. Each with name, company, role, trigger event candidate, and personalization line.
- Send the first 10 cold emails using the template in Section 1.9.

Day 4 (Thursday):

- Send the next 10 cold emails.
- Write the one-page posture summary template (specified in Section 2.6). Save as a fillable document ready for beta sessions.

Day 5 (Friday):

- Send the last 10 cold emails.
- Reach out to two network contacts who fit the ICP for a beta Signal Session. Use the beta pitch in Section 1.10.

Day 6 (Saturday):

- Write the LinkedIn headline and About section updates. Deploy to LinkedIn. The headline reads: "Fractional AI Risk Officer for Owner-Operator Firms | I help SMBs find out who is using AI with customer data, and fix it in 14 days | Based in Utah, working nationwide."
- Memorize the 11 high-signal session questions (Section 2.2). Run through them once aloud.

Day 7 (Sunday):

- Rest. Review Week 1 metrics: emails sent (30), beta contacts reached (2), site page live, intake form live, one-pager ready.
- Pre-write the Week 2 task sheet.

Day 8 (Monday):

- Follow up with Week 1 cold emails that did not respond. Use the follow-up template in Section 1.11.
- First beta Signal Session scheduled or in active conversation.

Day 9 (Tuesday):

- First beta Signal Session delivered. One-page posture summary produced within 24 hours.
- Send 10 more cold emails from new list additions.

Day 10 (Wednesday):

- Second beta Signal Session delivered. One-page posture summary produced within 24 hours.
- Extract lessons from both beta sessions and update the intake form or session question set if needed.

Day 11 (Thursday):

- Request testimonials from both beta clients using the template in Section 1.12.
- Send 10 more cold emails.

Day 12 (Friday):

- Any cold email reply in active conversation gets moved to the Calendly Signal Session booking flow.
- Send 10 more cold emails.

Day 13 (Saturday):

- Identify the two or three highest-probability prospects from the cold email thread. Send a trigger-event-specific follow-up (Section 1.13).
- Write the first LinkedIn post (one of the 12 misunderstood AI governance concepts in Section 2.9). Do not publish yet.

Day 14 (Sunday):

- Publish the first LinkedIn post.
- Review the first 14 days. Pipeline status: at minimum, two beta posture summaries delivered, two testimonials in hand or requested, 50 cold emails sent, one or two Signal Sessions booked in the pipeline.

1.3 The Revenue Targets

Cash commitments, not MRR abstractions.

- **Day 30:** First paid SHIELD Assessment invoiced. $3,500 collected (or$2,500 at stated launch rate if needed for first client only).
- **Day 60:** Second SHIELD Assessment invoiced at full $3,500. Running total:$7,000.
- **Day 90:** Third SHIELD Assessment or first Governance Sprint closed. Running total: $12,000 (three Assessments at launch rate plus full rate blend) to$16,500 (two Assessments at full rate plus one Sprint).

If the Day 30 target is missed by more than seven days, stop all content and infrastructure work. Focus exclusively on cold outreach and beta delivery calibration until the first paid Assessment is booked.

1.4 The Daily Non-Negotiables

Three inputs per day, every day, Monday through Friday.

1. **Outreach.** A minimum of five cold emails OR three cold LinkedIn DMs sent. Weekly floor: 20 emails or 15 DMs. Replies and follow-ups do not count toward this number.
2. **Delivery readiness.** At least 30 minutes practicing the 11 session questions aloud OR reviewing one section of Layer 3 (Practitioner Toolkit).

3. **Pipeline review.** Five minutes at end of day updating the prospect tracker. Who moved, who stalled, who got the next follow-up.

Content shipping does not count as outreach. Building does not count as outreach. Reading the playbook does not count as outreach. Only messages sent to a specific human asking for a specific conversation count as outreach.

This is the guardrail against the build-instead-of-outreach pattern that cost 30 cold emails with zero replies in March 2026. If the daily outreach floor is missed for more than two days running, stop all other work.

1.5 The Weekly Cadence

Every Friday, 30 minutes:

- Outreach volume sent vs target (minimum 20 emails per week)
- Replies received, meetings booked
- Pipeline movement (prospect stage changes)
- One thing to change next week

Every last Friday of the month, 90 minutes:

- Full funnel metrics (cold emails to replies, replies to Signal Sessions, Sessions to Assessments)
- Revenue vs target
- Methodology refinement based on delivery experience
- Decision log for the month

1.6 The Competitive Positioning Summary

When a prospect asks "how are you different from Faye?" or "why not a vCISO?" the answer is three sentences, memorized:

"Faye targets mid-market and enterprise. I work exclusively with owner-operator firms between 5 and 500 employees. A vCISO handles cybersecurity broadly. I focus specifically on AI risk, and I come from an operator background, not a cybersecurity one, which means I translate governance findings into margin and operational recommendations your team can act on without hiring anyone new."

Memorize. Deliver in under 20 seconds. Do not elaborate unless asked.

1.7 The List Construction Method

The first 30 qualified prospect names come from four sources, in order of quality.

1. **Utah-based owner-operator firms in professional services (10 names).** Source: Utah Chamber of Commerce member directories, Utah CPA Society member lists, ALA Utah (association of legal administrators), Utah Architects Association. Filter: 5-200 employees, owner-operator-led, professional services.

2. **Trigger-event-driven prospects (10 names).** Source: recent press coverage in Utah Business, Salt Lake Tribune business section, Inc. and Entrepreneur AI-incident coverage. Filter: any SMB mentioned in context of AI adoption, AI risk, vendor renewal, client due diligence, or cyber insurance.

3. **Boubacar's existing network (5 names).** Source: LinkedIn connections who run owner-operator firms in the target range. Already know Boubacar. Warm list, highest reply probability.

4. **Adjacent professional service firm owners (5 names).** Source: LinkedIn search with filter "CEO OR Founder OR Managing Partner" plus company size 5-200 plus location Utah OR Colorado OR Texas OR Illinois. Filter for owner-operator-led firms in accounting, legal, consulting, architecture, engineering, healthcare administration.

Spreadsheet columns: Name, Company, Title, Employee Count Estimate, Industry, Trigger Event (if any), Personalization Line, Email Address, Email Sent Date, Reply Date, Stage.

30 names by end of Day 3 is the minimum. Boubacar does not send a single cold email before 30 names are in the spreadsheet. Volume without qualification is spam. Qualification without volume is procrastination.

1.8 The LinkedIn Profile Setup

Headline: "Fractional AI Risk Officer for Owner-Operator Firms | I help SMBs find out who is using AI with customer data, and fix it in 14 days | Based in Utah, working nationwide"

About section (first paragraph visible above the fold): "I help owner-operator firms between 5 and 500 employees find out exactly how AI is being used in their business, where their exposure is, and what to do about it. 14 days, fixed fee, written report, policy in hand. The Fractional AI Risk Officer is the fractional executive model applied to AI governance and risk, specifically built for firms that do not have a CISO, general counsel, or dedicated compliance function."

Featured section: Link to catalystworks.consulting/governance. Link to the AI Tool Inventory Worksheet gated download. Link to Calendly for Signal Session booking.

Services section: List AI Risk and Readiness Assessment ($3,500, 14days).ListGovernanceSprint($ 8,500, 6-8 weeks). List Fractional AI Risk Officer Retainer (contact for scoping).

1.9 The Cold Email Template (V3)

Subject lines to test (send in rotation, measure reply rate weekly):

- "Quick question on your AI tool stack"
- "Your CRM's AI is reading your client emails"
- "[Their state] AI law and your vendor contracts"

Body:

```
Hi [First name],

[One sentence of specific personalization based on recent news about their firm or
their public content. Not generic. Invest 90 seconds per email.]

Quick question, and feel free to ignore if not relevant: do you know how many AI tools
are currently touching your customer data?

Most firms between 5 and 200 employees have 12 to 20. Most owners can name 3. The
other 9 to 17 are either embedded in software they already use (your CRM, your
accounting, your HR) or being used directly by employees through consumer accounts.

I help owner-operator firms find out what's in that list and fix what needs fixing in
14 days. Fixed fee, written deliverable, a policy you can implement the day you get
it.

Not pitching. If this is a question you're already thinking about, reply and I'll send
you the 15-question self-assessment I use with prospects. Takes 20 minutes. Will tell
you if you have a problem worth spending time on.

Boubacar Barry
Fractional AI Risk Officer | Catalyst Works Consulting
catalystworks.consulting/governance
```

The signature includes the title, the firm name, and one URL only. No phone number. No banner. No quote.

1.10 The Beta Session Pitch (for Network Contacts)

"[Name], I'm testing a new diagnostic I'm calling the Signal Session. 45 minutes, structured, I ask 11 questions about how AI is being used in your firm. You get back a one-page written summary of where you are and what to work on. No charge. I'm looking for two people in my network to run it with before I open it publicly. Would you be open to it? Takes 45 minutes plus a pre-session form that runs about 15 minutes to fill out."

Two beta sessions by Day 10. Both produce testimonials by Day 14.

1.11 The Follow-Up Template (for Cold Emails That Did Not Reply)

Send seven days after the first email.

```
Hi [First name],

Circling back once. If the AI tool stack question isn't on your radar this month, no
worries at all. I'll stop here.

If it is and the timing just got lost, the 15-question self-assessment is still
available. Reply with one word ("yes") and I'll send it over.

Boubacar
```

One follow-up maximum. No email three.

1.12 The Testimonial Request Template (for Beta Clients)

Send within 48 hours of delivering the posture summary.

```
Hi [First name],

Thank you for doing the beta Signal Session with me. The posture summary should have
landed in your inbox by now.

Quick favor. I'm building out the AI governance practice at Catalyst Works and the
strongest thing I can do to get it in front of the right people is share what you took
away from the session.

Would you be willing to share a short testimonial? Two or three sentences is enough.
Something like: what you thought before, what you found out in the session, what
you'll do next. I'll send you a short draft to react to if that's easier than writing
from scratch.

I'll use it on catalystworks.consulting/governance and on LinkedIn. Always anonymized
if you prefer (I'll just say "managing partner at a professional services firm in
Utah" or similar).

Whichever format works, thank you.

Boubacar
```

If the response is positive, send a pre-drafted two-sentence testimonial in the reply for the client to edit. Pre-drafting reduces friction and increases response rate.

1.13 The Trigger-Event Follow-Up (Send Selectively)

For the two or three highest-probability prospects from the cold email thread who did not reply to the original email or the Day 7 follow-up. Send only when there is a real trigger event.

```
Hi [First name],

[Name the trigger event specifically. Example: "Colorado just confirmed the June 30
enforcement date for SB 24-205 AI Act." Or: "[Their industry] just had a public AI
incident at [Company]."]

The 15-question self-assessment I mentioned in March is now pinned to that kind of
event specifically. If the [regulatory deadline / industry incident / vendor change]
is relevant to your business, reply with one word and I'll send it over. No call, just
the form.

Boubacar
```

Max two of these per month per prospect. Manufactured urgency kills reply rate.

1.14 The Week 1 Success Criteria

By end of Day 7:

- SHIELD one-pager written and saved
- Intake form built and linked
- catalystworks.consulting/governance updated with $3,500 pricing
- 30 qualified prospect names in the spreadsheet
- 30 cold emails sent
- Two beta contacts reached out to
- LinkedIn profile headline and About section updated
- 11 session questions memorized

By end of Day 14:

- Two beta Signal Sessions delivered
- Two posture summaries sent
- Two testimonial requests sent
- First LinkedIn post published
- 50+ cold emails sent cumulative
- At least one live conversation in pipeline toward a paid Assessment

If Day 14 arrives with no live conversation in pipeline, the issue is one of three things: the cold email copy is not landing (test a new subject line), the list is not qualified (add more Utah-based owner-operators), or the outreach volume is too low (double the weekly target to 40 emails). Diagnose by metric, not by feeling.

1.15 The Week 1 Kill Switches

Stop and reassess if any of the following are true by Day 14:

- Zero cold email replies of any kind (copy or list is broken)
- Both beta sessions produced no testimonial and no enthusiasm (methodology or delivery is broken)
- More than 20 hours spent this week on infrastructure vs outreach (the build-instead-of-outreach pattern is recurring)

The kill switch is not "stop doing the work." It is "stop the current approach, diagnose the specific failure, and fix one thing before continuing."

Layer 2: The Strategy

The operating strategy for the practice. Consulted when a strategic question arises. Not read daily.

Executive Summary

This document is the operating strategy for Catalyst Works Consulting's AI Governance practice. It is not a white paper. It is a working playbook.

The positioning: Boubacar Barry operates as a Fractional AI Risk Officer for owner-operator professional services firms with 5 to 200 employees. The offer leads with AI governance as the foot-in-the-door engagement, priced and scoped to generate first revenue while creating the structural conditions for follow-on margin and operational work.

The conversion thesis: every AI governance engagement ends with a written Margin and Operational Risk Indicators briefing that names the bottleneck and margin issues the AI usage is masking. This briefing is the bridge to the deeper Catalyst Works practice. Governance is not a loss leader. It is a fixed-fee, fixed-scope engagement that generates revenue on its own and makes the next conversation inevitable.

The category claim: Fractional AI Risk Officer. A specific, ownable role that builds on the fractional executive model SMB owners already understand (fractional CFO, fractional CMO). Specific enough to defend against generalist AI consultants. Wide enough to expand into retainer work without rebranding.

The methodology: SHIELD. Six steps, memorable acronym, defensible as Catalyst Works-owned IP.

Section 1: The Landscape

1.1 The Governance Triangle

Three overlapping disciplines that clients conflate constantly. Clarity here signals expertise fast.

- **IT Governance**: The factory. Infrastructure, availability, cybersecurity, technology-to-strategy alignment. Keeps the systems running and safe from outside attack.

- **Data Governance**: The raw materials. Data quality, lineage, access controls, lifecycle. Ensures data is accurate and usable.
- **AI Governance**: The finished product. The unique behaviors and risks of algorithms. Ethics, bias, explainability, accountability, human oversight. Governs how decisions get made by non-human agents.

The core truth: you cannot have credible AI governance without reliable data governance. You cannot have reliable data governance without functional IT governance. They stack.

1.2 The Six Core Principles in Practice

1. **Accountability**: Every deployed AI system has a named human owner.
2. **Transparency**: Users know when they are interacting with AI or AI-generated output.
3. **Fairness**: Models are tested for discriminatory bias before launch.
4. **Safety and Security**: Protection against adversarial attacks and against shadow AI (employees uploading sensitive data to unapproved tools).
5. **Explainability**: Constrained, rule-based AI workflows are preferred over black-box decision models for high-stakes actions.
6. **Human Oversight**: Workflows are designed so AI drafts, humans approve.

1.3 The Legal and Framework Map (verified as of April 2026)

The Utah AI Policy Act (SB 149). Effective May 1, 2024. The first US state law regulating private-sector use of generative AI. Applies to any individual or entity using generative AI to interact with people in Utah, regardless of where the organization is located. Administrative fines up to $2,500 per violation. For Catalyst Works specifically, this is the home-state authority play. Boubacar lives in Utah. The Utah AI Act has been in effect for nearly two years. Most Utah SMBs have not yet formally addressed it.

The Texas Responsible AI Governance Act (TRAIGA / HB 149). Effective January 1, 2026. Intent-based liability framework for AI misuse (unlawful discrimination, manipulation, deepfakes, child exploitation). Exclusive enforcement by the Texas Attorney General. Civil penalties up to $200,000 per violation with a 60-day cure period. Applies to any entity conducting business in Texas or offering products to Texas residents.

Illinois HB 3773. Effective January 1, 2026. Amends the Illinois Human Rights Act. Prohibits employers from using AI in ways that discriminate against protected classes. Requires employer disclosure when AI is used in recruitment, hiring, promotion, discipline, or termination decisions. Applies to employers with one or more employees in Illinois.

The Colorado AI Act (SB 24-205). Effective date: June 30, 2026 (after the delay from February 1, 2026 via SB 25B-004). Requires developers and deployers of high-risk AI systems to use reasonable care to protect against algorithmic discrimination. Subject to active political revision; a repeal-and-replace bill was under consideration in early 2026. Colorado stakeholders should monitor state legislative action through mid-2026.

The EU AI Act. In force since August 1, 2024. Banned practices enforceable since February 2, 2025. GPAI obligations since August 2, 2025. High-risk system requirements scheduled for August 2, 2026, though AI Omnibus negotiations may push this to December 2027 or August 2028. Fines up to €35 million or 7% of worldwide annual turnover (whichever is higher) for prohibited practices; up to €15 million or 3% for high-risk violations; up to €7.5 million or 1% for incorrect information.

NIST AI Risk Management Framework. US federal voluntary guidance. Four functions: Govern, Map, Measure, Manage. Not mandatory but increasingly referenced in litigation, procurement due diligence, and regulatory interpretation. Voluntary is a legal status, not a recommendation to skip it.

ISO/IEC 42001. International certifiable AI management system standard. Published 2023. Emerging as the "certification" version of AI governance for organizations that need a third-party verifiable claim. Catalyst Works does not claim 42001 expertise in year one; positions 42001 as a potential roadmap destination for Tier 3 clients.

Federal preemption risk. A December 11, 2025 executive order directed the DOJ to challenge state AI laws. Federal preemption typically requires congressional action. The legal landscape remains unstable. Messaging should anchor to operating-environment pressures (vendor questionnaires, client due diligence, insurance carrier requirements) rather than to specific regulatory deadlines that may move.

1.4 How SMBs Actually Implement Governance (and fail to)

The dominant implementation pattern in mid-market and enterprise is "top-down organizationally, bottom-up architecturally." Leadership sets the Acceptable Use Policy. Engineering implements an AI Gateway (a secure proxy that enforces rules technically).

For SMB owner-operators with 5 to 200 employees, the AI Gateway implementation is out of reach. No dedicated IT. No engineering team. The "top-down organizationally" part is the only lever available, and it has to be made to work without technical enforcement. This is the core design constraint for any SMB governance program and the reason enterprise-coded approaches fail in the SMB market.

Common failure modes:

1. **Wait-and-see paralysis.** Banning AI outright leads employees to use unsanctioned personal tools. Shadow AI is infinitely riskier than governed AI because it leaves no documentation trail, no training opportunity, no incident response path.
2. **Vendor blindspot.** Owners vet the primary LLMs (ChatGPT, Claude) while ignoring that their CRM, HR platform, accounting software, and project management tools have silently added AI features trained on customer data.
3. **Template theater.** Generic AI policies downloaded from the internet and stored on a shared drive. Signed once, never read again, not reflective of actual operations. Worse than no policy because it creates a paper record of compliance that fails under scrutiny.
4. **Decision-owner ambiguity.** No one is named accountable. When something goes wrong, the response is reactive and political rather than procedural.

1.5 The Anatomy of a Working SMB AI Policy

Seven sections, each one concrete enough that an employee can use it without asking a question.

1. **Purpose and Scope**: Why the policy exists and who it applies to.
2. **Risk Classification System**: Low / Medium / High / Prohibited, with one-sentence rules for each tier.
3. **Acceptable Use Rules**: What employees can use AI for, what they cannot.
4. **Data Handling and Privacy**: What data is explicitly prohibited from entering AI tools (client PII, financial data, protected health information, contract terms, etc.).
5. **Vendor Procurement Addendums**: The three clauses every AI vendor contract must include or disclose.
6. **Human Oversight Matrix**: What decisions require human review before AI output is delivered.
7. **Incident Response Protocol**: What happens when AI produces an error that affects a client.

Review triggers, not review calendars. The policy gets refreshed when a new regulation takes effect, a new AI vendor category is adopted, or an incident occurs. Not on a quarterly schedule that nobody remembers.

Section 2: The SHIELD Methodology

SHIELD is the Catalyst Works-owned methodology for AI Risk and Readiness Assessment. Six steps. Delivered in 14 days. $3,500 fixed.

- **S** (Surface): Complete AI tool inventory. Every tool touching the business, including embedded AI in software already in use.
- **H** (Halt): Identify the highest-risk tools and the highest-risk uses. Pause them where necessary. Document what's happening in the meantime.
- **I** (Inventory): Formal register with risk tier, accountable owner, data types, vendor terms summary.
- **E** (Educate): Employee communication that translates the policy into plain-language rules the team can actually follow.
- **L** (Lock): Written Acceptable Use Policy and vendor approval process, customized to the client's industry and jurisdiction.
- **D** (Document): Final deliverable with Governance Maturity Briefing and Margin and Operational Risk Indicators section.

The SHIELD deliverable is the foot-in-the-door product. The Margin and Operational Risk Indicators section is the bridge.

Section 3: The Offer Stack

Four productized offers. Each has a clear next step at 2-3x the price. No cliffs.

3.1 The Signal Session ($497)

Forty-five minutes. Structured diagnostic conversation via video call. Delivered with a one-page written posture summary within 24 hours of the session.

What happens: Boubacar runs 11 high-signal questions from the full 26-question diagnostic (Section 2.2 in Layer 3). The session maps the prospect to one of the four maturity levels. The posture summary names the level, the two or three highest-risk findings, and a specific recommended next step.

Conversion point: The posture summary recommends the SHIELD Assessment for any Level 0 or Level 1 client. For Level 2 clients, it recommends a targeted Governance Sprint. For Level 3 clients, it recommends a Fractional AI Risk Officer quarterly review retainer.

Intake process: Prospect books via Calendly. Confirmation email includes a 15-question intake form (pre-selected from the 26-question diagnostic, specified in Section 2.3 of Layer 3). Boubacar reviews the intake before the session and tailors the 11 session questions accordingly.

Why $497 and not free: A priced diagnostic signals that time is valuable and filters out tire-kickers. Boubacar's time is the constraint. The Signal Session price is a qualification gate

disguised as a service.

3.2 The SHIELD Assessment ($3,500, 14 days)

The core productized offer. Fixed price, fixed timeline, fixed deliverable set.

What clients buy: Fourteen days to find out exactly how AI is being used in their business, where the exposure is, and what to do about it. Written report. Policy ready to implement.

The six SHIELD phases (internal methodology):

- **S** (Surface): Identify every AI tool in active use, including shadow AI by individual employees.
- **H** (Halt): Flag the two or three highest-risk uses for immediate remediation. Not a legal opinion; a risk triage.
- **I** (Inventory): Document all AI tools, data inputs, output types, and user populations.
- **E** (Educate): Plain-language briefing to the leadership team on current exposure relative to applicable regulations.
- **L** (Lock): Draft the minimum viable AI Acceptable Use Policy customized to the client's industry and jurisdiction.
- **D** (Document): Full governance report with maturity level, findings, and 90-day action plan.

Deliverables:

1. AI tool inventory (full)
2. Risk triage summary (executive-readable, two pages)
3. Minimum viable AI Acceptable Use Policy (editable, signed off by Boubacar)
4. Governance report with maturity level and 90-day action plan
5. Margin and Operational Risk Indicators section (the bridge to follow-on work)
6. 30-minute Strategic Briefing call, delivered 7 days after final report

Disqualification criteria (applied at Qualification Call, before accepting any engagement):

- Fewer than 5 employees (insufficient AI exposure)
- More than 500 employees (scope requires enterprise team, not solo practitioner)
- No current AI usage (nothing to govern yet; refer to lead magnet and educational content)
- Needs a legal opinion on AI liability (refer to counsel)
- Already has a formal compliance program managed by a dedicated team

Launch rate (internal only, never published):
$2,500 for the first two engagements. Available as a stated discount to early clients. Pu$

3,500. After the first two engagements, the launch rate is retired.

Payment terms: 50% on contract signature, 50% on delivery of final report. Wire or Stripe.

Refund policy: Full refund if Catalyst Works cancels before kickoff. After kickoff, no refunds. If the client is dissatisfied with the deliverable, Catalyst Works revises once at no additional charge.

3.3 The Governance Sprint ($8,500, 6-8 weeks)

Implementation engagement. Operationalizes the recommendations from the SHIELD 90-day action plan.

Scope: Policy implementation, staff training framework, vendor AI usage audit, incident response protocol, and quarterly review cadence setup.

How it converts: Every SHIELD Assessment ends with a 90-day action plan. The Sprint is the offer that implements it. The conversion language is written into every SHIELD deliverable: "This action plan is designed to be self-implemented. If you would prefer a supported implementation, the Governance Sprint covers all items in 6-8 weeks at $8,500."

3.4 The Fractional AI Risk Officer Retainer ($1,500 to 5,000 per month)

Ongoing AI governance oversight. Six-month minimum engagement.

Tier 1 Essential ($1,500/month):

- Monthly governance review call (45 minutes)
- Policy updates triggered by regulatory changes
- New tool review on request (up to two per month)
- Incident triage availability (business hours)

Tier 2 Standard ($3,000/month):

- Everything in Tier 1
- Quarterly full AI inventory refresh
- Staff Q&A session (one per quarter)
- One additional ad hoc assessment per quarter

Tier 3 Full ($5,000/month):

- Everything in Tier 2
- Dedicated Slack or Teams channel

- Custom-configured governance monitoring agent (deferred to post first retainer signed)
- Attendance at one internal leadership meeting per quarter

How it converts: The Sprint naturally produces retainer conversations because the client has now invested $12,000 in governance and does not want the program to decay. Every Sprint engagement ends with a retainer proposal.

3.5 Pricing Logic

$497 Signal Session \rightarrow$3,500 SHIELD Assessment → $8,500 Sprint OR$1,500 to $5,000/month Retainer.

Every offer has a clear next step at 2-3x the price. No gaps, no cliffs. The funnel compounds because the prospect always sees what comes next.

Section 4: Marketing and Go-to-Market

4.1 Web Architecture

Decision: AI governance content lives as a dedicated section on catalystworks.consulting at /governance. No separate microsite. Rationale: SEO concentration on a single domain, brand unity, maintenance overhead for a solo operator, and exposure of governance traffic to bottleneck and margin content (the conversion thesis).

4.2 YouTube Series

Title: *AI Without the Lawyer: Field Notes for Owner-Operators*

Cadence: bi-weekly. 8-12 minutes per episode. Talking head with selective screen-share.

Season 1 structure: 8 episodes, outlined separately. Episode 1 script outline included as a working artifact.

4.3 Seminars

Monthly, 45 minutes plus Q&A, live with replay for attendees only. First three sessions outlined in working artifacts. Seminar-to-Assessment conversion target: 3-4% of registrations convert to paid Assessments within 60 days.

4.4 LinkedIn Organic

3-5 posts per week. Format mix: 40% short-form tactical text, 30% personal observation, 20% jurisdiction-specific explainer, 10% YouTube clip or seminar excerpt. Tuesday-Thursday posting at 7:30am MT.

4.5 Cold Outreach

Two parallel sequences: margin-led (existing subject line: "Where is your margin actually going?") and governance-led (new subject: "Quick question on your AI tool stack"). 20 emails per sequence per week. Reply rates tracked separately. Volume rebalanced after 60 days based on what converts.

4.6 Podcast and Speaking

From week 3 onward, 2 podcast pitches per week. Targets: SMB-focused podcasts, state CPA society podcasts, industry-vertical podcasts for professional services. Goal: first guest appearance recorded by week 10, published by week 12.

4.7 Content Repurposing System

Every YouTube episode → 4 LinkedIn posts. Every seminar → 5 derivative pieces (YouTube edit, LinkedIn clip series, newsletter article, slide deck, audio cut). Every blog post → LinkedIn teaser with excerpt.

Section 5: The Conversion Thesis

5.1 Governance Findings That Surface Margin Issues

This table is the practitioner's primary translation tool. Every SHIELD Assessment surfaces at least two entries from this table in the Margin and Operational Risk Indicators section. Each finding has a client-facing translation that makes the cost of inaction tangible.

Governance Finding	Bottleneck or Margin Equivalent	Client-Facing Translation
Shadow AI usage by employees	Decision-making bottleneck (no approved tools, workers default to unauthorized solutions)	Your team is solving problems faster than your processes allow, without guardrails
No AI acceptable use policy	Accountability gap (no owner, no response protocol)	If something goes wrong with an AI tool today, no one knows who is responsible
Vendor AI policies unreviewed	Vendor risk concentration (third-party AI embedded in core operations)	Your vendors may be training their models on your business data right now
AI outputs in client work without disclosure	Client trust and compliance risk	A client who finds out you used AI without telling them is a client who does not renew
No AI incident response plan	Operational fragility	One bad output sent to the wrong person and there is no playbook for what happens next
Training data includes customer PII	Regulatory exposure (state and federal AI laws)	Customer data in AI training pipelines is the most common path to regulatory action in 2026
No maturity baseline established	Invisible progress (no measurement, no improvement signal)	You cannot get better at something you have never measured
AI tool proliferation without procurement review	Budget fragmentation and security surface expansion	You are paying for tools that duplicate each other and creating entry points you cannot see

5.2 The Strategic Briefing Structure

Delivered 7 days after the final SHIELD Assessment report. 30 minutes. Included in the Assessment fee.

Agenda:

1. Walk through the Margin and Operational Risk Indicators section of the deliverable (10 minutes)
2. Connect the three most significant indicators to specific next-engagement options (10 minutes)
3. Propose the two models for continued work: Governance Sprint ($8,500, fixed) or Fractional AI Risk Officer Retainer ($ 1,500 to $5,000/month) (5 minutes)
4. Answer questions, establish next steps, send engagement letter within 48 hours (5 minutes)

5.3 Disqualification Criteria

Applied at the 15-minute Qualification Call, before any paid engagement. Repeated here because disqualifying the wrong buyer is the single highest-leverage move a solo operator can make.

- Under 5 employees: refer to lead magnet and seminar; not ready for $3,500 Assessment.
- Over 500 employees: outside ICP; refer or pass. Borderline cases (200-500) require scoping before pricing.
- No AI usage and no adoption plan: not the buyer for this offer.
- Requires legal opinion as a deliverable: not the offer; refer to counsel.
- Existing governance program in place: gap-analysis Assessment at $5,500 instead of $ 3,500, given additional complexity.

5.4 The Practice-Level Failure Modes

Six named failure modes for the practice itself. Monitored monthly.

1. **Selling to the wrong buyer.** The wrong buyer wants governance as a checkbox, not as risk reduction. They underinvest in implementation and produce weak or no testimonials. Mitigation: the intake form question "What is driving your interest in AI governance right now?" filters motivated buyers from checkbox buyers.

2. **Underpricing the diagnostic.** Dropping below $2,500 for any SHIELD Assessment signals the methodology is negotiable. Every future conversation then starts with a discount request. Mitigation: state the price once, clearly, without qualification. No discounts after the first two engagements.

3. **Skipping the disqualification step.** Taking an engagement that fails one or more disqualification criteria creates scope creep and delivery failure. Mitigation: apply disqualification criteria in the first five minutes of any sales conversation, before methodology or pricing.

4. **Treating the Assessment as the end state.** Terminal engagements leave $17,000 or more on the table per client over 12 months. Mitigation: the conversion path is written into every deliverable. The 90-day action plan always includes the Sprint offer. The Sprint always includes the retainer proposal.

5. **Content before clients.** The content engine is right in sequence and format but belongs in weeks 3 through 52, not weeks 1 or 2. Mitigation: the first LinkedIn post ships in week 3, after the first beta session.

6. **Tool-building before revenue.** The governance agent, monitoring dashboard, and Supabase schema are all deferred to post first revenue. Mitigation: the governance agent build is gated on $12,000 in cumulative Assessment revenue. No exceptions.

Section 6: Execution Cadence

6.1 The Forge 2.0 Daily and Weekly Rules

Non-negotiable weekly outreach quota:

- 20 governance cold emails sent by Monday noon
- 5 LinkedIn DMs sent by Wednesday
- 2 podcast pitches sent by Friday

Content shipping does not count as outreach. Building does not count as outreach. The only thing that counts as outreach is a message sent to a specific human asking for a specific conversation.

This is the explicit guardrail against the build-instead-of-outreach pattern that cost 30 cold emails with zero replies in March 2026.

6.2 Weekly Review

Every Friday, 30 minutes:

- Outreach volume sent vs target
- Replies received, meetings booked
- Content shipped vs planned
- Pipeline movement (prospect stage changes)
- One thing to change next week

6.3 Monthly Review

Last Friday of each month, 90 minutes:

- Full funnel metrics (leads, qualification calls, assessments, conversions)
- Revenue vs target
- Content performance (YouTube views, seminar attendance, lead magnet downloads)
- Methodology refinement based on client feedback
- Decision log for the month

Section 7: Mastery and Progression

7.1 Must Master Before First Paid Engagement

- SHIELD methodology, articulated in 60 seconds without notes
- Full 60-minute discovery session flow, practiced twice in beta
- Maturity Model scoring, applied live in session
- Objection responses for all seven common objections, spoken naturally
- Regulatory landscape summary for Utah, Texas, Illinois, Colorado, EU in plain language

7.2 Develop Through First 3-6 Months

- Industry-specific risk patterns (accounting, legal, consulting, agency, architecture, engineering)
- Vendor contract reading (read three in full: OpenAI, Microsoft, Salesforce)
- Policy language craft (write for specificity, not generality)
- Strategic Briefing delivery (the conversion-critical 30-minute call)
- Content flywheel management (sustaining output while delivering engagements)

7.3 Long-Term Depth (Unlocked Only After Tier 2 Milestones Hit)

These items are not started until $10K MRR is sustained and the first retainer client has renewed. Listed here only for forward visibility.

- ISO 42001 working knowledge (not certification, yet)
- Sector-specific governance expertise in one or two verticals
- Practitioner community standing (IAPP presence, conference speaking)

- Named methodology evolution (SHIELD v2, licensable version as SHIELD Certified Practitioner program)

Section 8: The Deferred List

These items were in earlier versions of the strategy. They are deferred to specific unlock criteria, not abandoned. Starting any of them before the unlock triggers is the build-instead-of-outreach pattern.

- **The governance agent (bot-driven diagnostic on the website).** Unlock: $12,000 cumulative Assessment revenue collected.
- **The YouTube series beyond Episode 1.** Unlock: first paid Assessment invoiced. Episode 2 ships Week 10 at earliest.
- **Monthly free seminars.** Unlock: warm email list above 200 subscribers. Deferred from original Week 6 target.
- **Lead magnets 2 and 3 (Compliance Map, Vendor Decoder).** Unlock: Lead Magnet 1 has been downloaded more than 50 times. Defer to Week 9+ earliest.
- **Dedicated microsite at AIGovSMB.com or equivalent.** Unlock: /governance traffic above 5,000 monthly visits.
- **ISO 42001 Lead Implementer certification.** Unlock: month 12-18, evaluate based on client demand.
- **Multi-jurisdictional EU AI Act compliance advisory work.** Unlock: year 2, outside current ICP until then.
- **The 14-section governance policy full build.** Current SHIELD offer uses the 7-section policy; 14-section version is a Tier 3 upsell, not the standard deliverable.
- **Productized book or course (SHIELD Method).** Unlock: $25K cumulative cash collected and three completed Assessments with testimonials.
- **SHIELD Certified Practitioner licensing program.** Unlock: three-client base stable at retainer, month 18 minimum.
- **The Catalyst Council Review offer.** Unlock: year 2 minimum. Not in the primary offer stack.
- **Scaling beyond solo operator (hiring, subcontracting, white-label).** Unlock: retainer base stable at $10K MRR minimum for three consecutive months.

Section 9: Success Criteria

30-Day Cash Target

- First paid SHIELD Assessment invoiced and collected: $3,500(or$2,500 at stated launch rate for client 1)
- Two beta Signal Sessions delivered with posture summaries
- Two testimonials captured
- /governance section live with SHIELD Assessment publicly priced at $3,500
- 60 cold emails sent cumulative
- First LinkedIn post published (Week 3)

60-Day Cash Target

- Second SHIELD Assessment invoiced at full $3,500
- Running total: $7,000
- One Governance Sprint conversation in active proposal or scoping
- First lead magnet (AI Tool Inventory Worksheet) published and gated

90-Day Cash Target

- Third SHIELD Assessment invoiced OR first Governance Sprint closed
- Running total: $12,000to$16,500 depending on mix
- First retainer conversation at or near close
- Second lead magnet optionally published
- First YouTube episode shipped (deferred from original Week 4 plan; now Week 10+)

6-Month Checkpoint

- $10Kto$15K cumulative cash collected
- One active retainer ($1,500to$5,000/month)
- Two to three additional SHIELD Assessments delivered
- First podcast guest appearance booked or recorded
- Book outline started only if the above milestones are met

Appendix: Deliverable Checklist

The working artifacts that accompany this strategy document:

1. Productized Offer One-Pager, AI Risk & Readiness Assessment
2. Episode 1 YouTube Script Outline, The Five AI Tools Already Touching Your Client Data
3. Lead Magnet 1 Outline, AI Tool Inventory Worksheet for Owner-Operator Firms
4. First Seminar Outline, AI Without the Lawyer (45-minute walkthrough)
5. Revised Cold Email for Governance Angle
6. Web Property Recommendation Memo
7. 12-Week Roadmap (structured table)
8. Rebuilt AI Governance Practitioner Toolkit (see Part II below)
9. This strategy document (Part I)

End of Part I. Part II (Practitioner Toolkit) follows.

Layer 3: The Practitioner Toolkit

Internal-facing working reference used during client engagements. Not for client distribution.

The material below is internal-facing working reference, used during client engagements. Not for client distribution.

Toolkit Section 1: The Discovery Session

1.1 Session Architecture

This is a 60-minute structured working session, not a sales call. The frame matters: you are conducting a diagnostic, not pitching a service.

Call it what it is: an AI Governance Diagnostic Session.

Phase	Time	What Happens
Open	0-5 min	Frame the session. Set the agenda. Establish that you are here to understand their situation, not to sell them anything today. One sentence on what you do, then stop talking about yourself.
Current State	5-20 min	Diagnostic questions: AI usage inventory, risk exposure, existing policies. Listen more than you speak. Take visible notes. Surface the gap between what they think they have and what they actually have.
Accountability Map	20-30 min	Identify who owns what. Who is responsible for AI decisions today? Who gets called if something goes wrong? What happens when an employee uses an unapproved tool? This is where most SMBs discover they have no answers.
Vendor Exposure	30-40 min	Walk through their AI vendor stack. For each tool: what data flows into it? What does the vendor's contract say about retention and training? Most clients have not read these clauses.
Maturity Score	40-50 min	Score them live on the maturity model (Section 1.3). Do this visibly. Live scoring creates a shared reference point and signals measurement, not improvisation.
Closing Move	50-60 min	Reflect back what you heard. Name the top two or three gaps. Present the output template (Section 1.4) and explain it will arrive within 48 hours. Close with one specific question: "Is fixing this in your top five priorities right now?"

1.2 Opening Framing Script

"Today I want to understand how your organization is actually using AI, who is accountable for those decisions, and where your exposure is. I am going to ask specific questions and take notes. At the end, I will give you a score on a governance maturity scale and tell you the two or three things that need attention most urgently. You will have a written summary within 48 hours. Sound useful?"

1.3 Closing Move (The One Question)

"Based on what we just found, [name the top gap specifically: 'you have no approved vendor list and three tools with client data in them'] is closing that kind of exposure in your top priorities right now? Because if it is, I have a specific engagement that addresses exactly this."

1.4 Diagnostic Questions by Category

These are the exact questions. Do not ask all of them. Use the first one in each category to gauge depth, then follow the thread. Your job is to surface the gap, not exhaust the topic.

CATEGORY A: Current AI Usage

- A1: Walk me through every AI tool your team uses on a regular basis, including tools built into software you already use (AI features in your CRM, your email, your accounting software).
- A2: Are there tools your employees are using informally, outside of anything you have officially approved?
- A3: Which of these tools touches client data, employee data, or proprietary business information?
- A4: Who made the decision to adopt each of these tools, and what was the approval process?
- A5: Is there an AI system that is making or significantly influencing decisions that used to be made by a human?

CATEGORY B: Risk Exposure

- B1: If one of your AI tools gave a client wrong information or made a discriminatory decision, what would happen? Who would handle it, and what is your liability exposure?
- B2: Have you had any situation where an AI tool produced an output that surprised you, embarrassed you, or caused a problem?
- B3: If a client asked you today how you use AI and what data protections you have in place, what would you say?
- B4: Are any of your clients in regulated industries (healthcare, financial services, legal, HR)? Do you know if your AI use creates compliance obligations for them?
- B5: What would a data breach involving an AI vendor cost you financially, reputationally, in client relationships?

CATEGORY C: Existing Policies

- C1: Do you have a written policy that governs how your team uses AI? If yes, where does it live and when was it last updated?
- C2: Does your employee handbook say anything about AI use?
- C3: What data are employees explicitly prohibited from inputting into AI tools? Is this written anywhere or understood informally?

- C4: If I asked one of your employees tomorrow what they are and are not allowed to do with AI, what would they say?
- C5: Have you reviewed your client contracts recently to check whether your AI use is consistent with your data handling obligations?

CATEGORY D: Stakeholder Accountability

- D1: If something goes wrong with an AI system tomorrow, who is the first call? Who owns the response?
- D2: Who in your organization has authority to approve or reject the adoption of a new AI tool?
- D3: Does your leadership team have a shared understanding of what your AI risk exposure is?
- D4: Has AI governance ever come up in a board meeting, an investor conversation, or a client due diligence process?
- D5: If a regulator or auditor asked you to demonstrate responsible AI use, could you produce evidence of that within 48 hours?

CATEGORY E: Vendor Relationships

- E1: Have you read the AI-specific data provisions in your vendor contracts? Specifically: does your vendor use your data to train their models?
- E2: Do you have a list of AI vendors your team is approved to use? Who maintains it?
- E3: When you evaluate a new AI software vendor, what specifically do you ask about their AI practices?
- E4: If a vendor's AI system failed or was compromised, what notification would you receive and how quickly?
- E5: Do you have any contractual protections (indemnification, limitation of liability) covering AI-related vendor failures?

1.5 Governance Maturity Model

Four levels. Score the client live during the session. Each level has a name, what it looks like in practice, the immediate risk it represents, and the specific next step.

Level	Name	What It Looks Like	Immediate Risk	Next Step
LEVEL 0	Ungoverned	No AI inventory. No written policies. Employees using AI tools at discretion. Leadership has no visibility.	Maximum. Client data in unapproved tools. No incident response capability. Regulatory exposure if any state AI law applies.	Complete AI inventory first. Then Acceptable Use Policy within 30 days.
LEVEL 1	Aware but Informal	Leadership knows AI is being used. Some informal norms exist. No written policy. No inventory. Vendor contracts unreviewed.	High. Informal norms are unenforceable and unevenly applied. No documentation means no defense if challenged.	Document the norms. Build an inventory. Formalize Acceptable Use Policy. Assign accountability.
LEVEL 2	Policy Exists	A written AI policy exists. Inventory partially complete. Accountability named for some systems. Basic vendor review process in place.	Moderate. Policy may not cover all use cases. Training may be absent. Incident response plan likely missing.	Complete inventory. Add incident response and human oversight protocols. Train staff. Establish review cadence.
LEVEL 3	Production-Ready	Complete inventory. Full policy with all required sections. Named accountable owners. Incident response tested. Staff trained. Quarterly review cadence.	Low. Governance is functional and defensible. Remaining risk is drift: policy aging faster than AI adoption.	Monitor and maintain. Consider ISO 42001 alignment as a trust signal. Use governance as a competitive differentiator.

1.6 Session Output Template

Delivered to the client within 48 hours. Designed to be forwarded internally, which is how it generates pull for follow-on work. Keep it to two pages maximum.

Deliver as a clean PDF with Catalyst Works branding. Title: "[Client Name] AI Governance Diagnostic | Catalyst Works Consulting | [Date]."

Section-by-section structure:

1. **What We Covered.** One paragraph, 3-4 sentences. Summarize session scope without editorializing. Establishes shared context for anyone not in the room.
2. **Your Governance Maturity Score.** State the score: Level 0 / 1 / 2 / 3. One sentence on what that level means for this specific organization. Do not soften it. Do not dramatize. A score is measurement, not judgment.
3. **Top Three Gaps Identified.** Specific, not generic. "No written acceptable use policy" is acceptable. "Governance documentation needs attention" is not. Each gap: one sentence on why it matters, one sentence on what unaddressed risk it creates.
4. **AI Tool Inventory (As Documented).** List every AI tool identified during session. Tool name / primary use / data types input / risk tier (L/M/H/Prohibited). Even an incomplete inventory makes visible what was invisible.
5. **Accountability Map (As Identified).** Who was named accountable for what. Gaps noted explicitly: "No accountable owner identified for [tool]."
6. **Margin and Operational Risk Indicators.** Three to five observations that connect governance gaps to operational or margin issues. This section is the conversion bridge. This is why Catalyst Works engagements are different from template-driven governance work.
7. **Recommended Next Step.** One specific recommendation tied to the maturity level. States what the outcome looks like and roughly how long it should take. This is where the follow-on engagement is introduced, not pitched.

Toolkit Section 2: Objection Handling

Do not rehearse these as scripts. Internalize the logic. The goal is genuine persuasion, not clever comebacks. Each response reframes the objection without dismissing it.

2.1 "We're too small to need this."

The NYC MyCity chatbot that gave small business owners advice to withhold workers' tips and ignore other labor laws was not an enterprise problem. Size determines the complexity of governance, not whether you need it. A three-person firm with one AI tool touching client data

is exposed. The scale of the governance matches the scale of the operation, which is exactly what I help design.

2.2 "We don't really use AI that much."

Does your CRM have AI features? Does your email platform? Your accounting software? AI is embedded in most business software now, and you may be using it without knowing it. The question is not whether you use AI. It's whether you know how it's using your data.

2.3 "We already have a privacy policy."

A privacy policy covers how you handle data you collect. It does not cover how your employees use AI tools, what data they put into those tools, how your vendors use that data, or who is accountable when an AI system makes a wrong call. Those are four separate governance gaps that a privacy policy does not touch.

2.4 "This feels like compliance overhead."

Here's the alternative: one of your employees puts a client contract into an AI tool whose vendor uses it to train their model. Your client finds out. You have no policy, no documentation, no accountability structure. The conversation you have then costs more (legal exposure, lost client trust, management time) than governance would have. This is not overhead. It's insurance that you pay once and use forever.

2.5 "We'll deal with this when regulations require it."

Texas TRAIGA went live January 1, 2026. Illinois HB 3773 went live the same day. Colorado's AI Act takes effect June 30, 2026. Utah's AI Policy Act has been live since May 1, 2024. If you have employees or clients in those states, it's already required. And the organizations that wait for a regulatory trigger always pay more to remediate than they would have spent to build it right the first time. Reactive governance is always more expensive than proactive governance.

2.6 "We don't have budget for this right now."

What's your legal bill if one of your vendors has an AI data incident and you're named in the client complaint? What does it cost you if a prospective client asks about your AI practices in a due diligence call and you have no answer? Governance is not a discretionary expense. It's a hedge against a specific class of operational and legal risk. What I do is find the minimum viable version that gives you that protection at a price that makes sense for your operation.

2.7 "Can't I just use a template I find online?"

You can. The same way you can use a contract template you find online. The risk is that a generic template does not reflect your specific AI tool stack, your actual data practices, your real vendor relationships, or the laws that apply to your specific jurisdiction and industry. And a governance policy that does not reflect your actual operations is not just useless. It's actively misleading if it ever gets scrutinized.

Toolkit Section 3: The Business Case for Skeptical Owner-Operators

Anchor every conversation to outcomes the client already cares about. Do not lead with compliance. Lead with what keeps them awake.

3.1 Liability

If an AI tool you use makes a wrong decision (misclassifies a client, gives bad advice, discriminates in hiring) and you have no governance documentation, you cannot demonstrate that you acted reasonably. Governance is what creates a paper trail of responsible practice. Without it, you are exposed. With it, you have a defense.

3.2 Operational Continuity

What happens when the AI tool your team depends on changes its pricing, gets acquired, or changes its data practices? A governance program includes vendor risk management. It tells you what tools you're dependent on, what your exit strategy is, and what the data implications of a vendor change are. Organizations without this find out the hard way.

3.3 Vendor Risk

Most SMBs have signed AI vendor contracts without reading the AI-specific clauses. Some vendors explicitly reserve the right to use your data to train their models. Some do not notify you of incidents unless required by law (which varies by jurisdiction). A governance program includes a vendor assessment process that asks the right questions before you're locked in.

3.4 Employee Behavior

Your employees are making decisions today about what data they put into AI tools. Without a written policy, those decisions are ungoverned. An employee sending client financials to an AI tool to generate a report is not malicious. They're solving a problem with the tools available to

them. Governance gives them a framework so they solve problems in ways that don't create new ones.

3.5 Competitive Positioning

Your enterprise clients and prospects are starting to ask about AI governance in due diligence. If you cannot answer "What is your AI governance policy?" you are at a disadvantage against competitors who can. Governance is increasingly a procurement requirement, not just an ethical nice-to-have.

Toolkit Section 4: Trigger Events

These are moments when a cold prospect becomes warm. Watch for them in news, client conversations, and social media.

- A competitor, industry peer, or well-known company in their sector has an AI-related incident (data breach, discriminatory output, chatbot failure). Adjacency is enough.
- A new state AI law takes effect in their operating jurisdiction. Regulatory deadlines create urgency positioning cannot manufacture.
- A client sends a vendor questionnaire with AI governance questions. Fastest-growing trigger. Enterprise procurement is adding AI to due diligence checklists.
- They're applying for a contract or grant that requires a compliance statement or evidence of responsible AI practices.
- A new employee, board member, or investor asks "what is your AI policy?" and they don't have an answer.
- They're onboarding a new AI tool and someone internally raises concerns about data security or liability.
- They experience an internal AI failure. This is the event that converts skeptics into buyers.

Signal monitoring: set Google Alerts for "AI governance lawsuit," "AI bias lawsuit," "AI data breach SMB," "AI vendor contract data." When these surface, that's a content hook and a prospecting context.

Toolkit Section 5: The Three-Tier Implementation Roadmap

Present this roadmap at the end of the diagnostic session or in the output template. It gives clients a destination and contextualizes the immediate next step within a larger journey.

5.1 Tier 1: Minimum Viable Governance ("Not Exposed")

Target: 1-50 employees. AI adoption is opportunistic. No dedicated IT or compliance function. Primary concern: not being caught flat-footed.

Deliverables:

- Complete AI inventory (every tool, every use, every data type)
- Acceptable Use Policy (what's permitted, what's prohibited, who approves new tools)
- Named accountable owner for each tool
- Basic vendor data assessment for highest-risk tools
- Employee communication (one-page summary of the policy)

Time to implement: 30-60 days with active client participation. 2-3 working sessions plus async review.

Common blockers: Client cannot find time to complete inventory. Leadership has not agreed on "prohibited use." Vendor contracts not accessible or unsigned PDFs in someone's inbox.

Signs ready for Tier 2: Inventory complete and current. Policy written, distributed, acknowledged. At least one accountable owner per tool. Client has adopted at least one new AI tool and used the approval process to do it.

5.2 Tier 2: Operational Governance ("Systematically Managed")

Target: 50-250 employees. AI embedded in multiple business functions. Some regulatory exposure or client due diligence requirements. Needs governance that works across teams.

Deliverables:

- Full multi-section governance policy
- Risk classification for all AI systems in inventory
- Human oversight protocols by risk tier
- Vendor management process with standard assessment checklist
- Incident response plan
- Role-specific employee training
- Quarterly review cadence established

Time to implement: 60-120 days. Cross-functional participation (legal, operations, HR, leadership). 4-6 working sessions plus document review.

Common blockers: Legal review delays. Disagreement on "high-risk" in this organization. Resistance from department heads who feel the policy constrains their teams. Training rollout competing with other organizational priorities.

Signs ready for Tier 3: Policy functional, not just documented. Staff can answer "what does our policy say about X" without looking it up. At least one incident has been processed through the response protocol. Client has used governance to evaluate and reject or approve a new vendor.

5.3 Tier 3: Mature Governance ("Differentiated")

Target: 250-500 employees, or high-stakes regulated industries at any size. AI is a strategic function. Governance is a competitive differentiator. May serve enterprise clients who require governance evidence.

Deliverables:

- ISO 42001 alignment (not necessarily certification)
- AI ethics statement for external communication
- Governance report for board or leadership
- Bias audit process for high-risk systems
- Governance-as-trust-signal communications strategy
- Vendor contract clause library
- Annual governance review as a formal event

Time to implement: Ongoing. Initial build 90-180 days. Governance is now a management function, not a project.

Common blockers: No one owns the ongoing governance function. Review cadence breaks down when leadership attention is elsewhere. Documentation drifts out of sync with actual adoption. ISO alignment requires organizational change beyond documentation.

Signs Tier 3 is working: Governance is referenced in sales materials and client conversations. New AI tool adoption is self-governed by staff using the established process. Leadership can answer "what is our AI governance maturity?" with a specific, evidence-backed answer.

Toolkit Section 6: Designing Governance That Stays Relevant

The failure mode is a policy written in 2026 that's still in version 1.0 in 2028 while the organization has adopted twelve new AI tools. These mechanisms prevent that.

- **Amendment triggers baked into the policy text** (not a separate process): specific events that automatically trigger review. New regulation. New AI vendor category. AI incident. Organizational restructure.
- **AI inventory review as a quarterly operations item:** tie it to an existing meeting (quarterly business review, operations review) so it doesn't depend on someone remembering.
- **Tool adoption as review trigger:** every time the organization adopts a new AI tool, the inventory and risk classification is the gate. No inventory update, no tool approval.
- **Sunset clauses on vendor assessments:** assessments expire after 12 months and must be renewed. Forces regular contact with vendor data practices.
- **Living document format:** policies stored in version-controlled documents (Notion, Google Docs with version history). PDFs are archives. Living documents are systems.
- **Regulatory monitoring:** build in a named person responsible for monitoring AI regulatory developments in the organization's jurisdictions. Can be a Google Alert or a paid monitoring service. Must be someone's job.

Toolkit Section 7: Content Angles and Educational Arc

7.1 The Most Misunderstood Concepts in AI Governance

Each one is a content piece waiting to exist.

- **"AI governance" and "AI ethics" are not the same.** Ethics is the values. Governance is the system that operationalizes values inside an organization. Ethics without governance is aspiration. Governance without ethical grounding is compliance theater.
- **Having a privacy policy does not mean you have AI governance.** A privacy policy governs how you handle data you collect from users. It does not govern how your employees use AI tools, what data they input, or how vendors use that data.
- **"Low risk" AI is not "no risk" AI.** Risk classification must account for scale and breadth of use, not just individual system risk level.
- **NIST AI RMF is voluntary, not optional.** Voluntary means no penalty for non-adoption. It does not mean no consequences. Organizations facing litigation, regulatory scrutiny, or client due diligence without a recognized framework are at a disadvantage.
- **"The vendor handles it" is not a governance answer.** Vendors handle their own data practices. They do not govern how your employees use the tool, what data they input, or who's accountable inside your organization for the AI system's outputs.
- **Governance is not a one-time project.** AI capabilities change quarterly. Regulations evolve monthly. A governance document written once and never reviewed is outdated by the time

the ink dries.

7.2 What SMBs Are Actually Searching For

These are real entry points for content.

- "Does [state] AI law apply to small businesses?" (Utah, Colorado, Texas, Illinois variants)
- "Do I need an AI policy for my company?"
- "Can my vendor use my data to train their AI?"
- "What is an AI acceptable use policy?"
- "ChatGPT data privacy business"
- "AI bias employment law"
- "What to put in an AI policy"
- "Is using AI with client data legal?"
- "AI vendor risk management"
- "How to govern AI in a small company"

7.3 Five Underserved Content Angles

1. **The AI Tool Audit Walkthrough.** Short video series plus companion checklist. Audience: owner-operators who suspect they have shadow AI but don't know how to find it. Gap: existing resources explain what shadow AI is; none walk through how to audit tool-by-tool, department-by-department.

2. **Vendor Contract AI Clause Decoder.** Template plus annotated example plus short video. Audience: business owners who have signed SaaS contracts and don't know what the AI provisions say. Gap: legal blogs explain AI clauses exist; no one has produced a plain-language decoder for SMBs.

3. **The "30-Day Minimum Viable Governance" Series.** Email sequence or YouTube series. Audience: small business owners who have just discovered they have no AI policy and need to fix it fast. Gap: governance resources are written for organizations with a compliance team; no one has produced a week-by-week sequence for a two-person operation.

4. **State AI Law Tracker for SMBs.** Regularly updated reference page plus email alerts. Audience: SMB owners operating in multiple states. Gap: legal firms publish summaries for enterprise legal teams; no one has produced an SMB-readable tracker mapping laws to business impact.

5. **The "What Happens When AI Goes Wrong" Case Study Series.** Written case studies (1,000-1,500 words each) plus short video summaries. Audience: owners not convinced the risk is

real. Gap: case studies exist for researchers and enterprise risk managers; none translate high-profile failures into specific lessons for SMB operators.

7.4 First-90-Days Content Arc

Sequenced by purpose. Early content builds credibility and positions vocabulary. Mid-sequence generates leads. Late-sequence creates infrastructure that compounds.

Week	Content Type	Topic
1-2	LinkedIn post	What AI governance is (and what it isn't). Three distinctions that clear up 80% of the confusion.
2-3	LinkedIn post	The statistic every SMB owner should know (verified statistic only; no hallucinated numbers).
3-4	YouTube Episode 1	The Five AI Tools Already Touching Your Client Data
4-5	Lead magnet	AI Tool Inventory Worksheet for Owner-Operator Firms
5-6	LinkedIn post	Can your vendor use your data to train their AI? Three clauses to find.
6-7	LinkedIn post	The four maturity levels of AI governance. Where does your organization sit?
7-8	LinkedIn newsletter	Utah's AI Policy Act has been in effect since May 2024. Most Utah SMBs have never addressed it.
8-9	Lead magnet	2026 SMB AI Compliance Map (Utah, Texas, Illinois, Colorado, EU)
9-10	LinkedIn post or video	Five AI tools your employees are probably using right now without governance coverage. Named specifically.
10-11	Case study	Anonymized: "How we built minimum viable governance for a 15-person professional services firm in 30 days."
11-12	Email sequence launch	30-Day Minimum Viable Governance email series

Toolkit Section 8: Mastery Tiering

8.1 Must Master Before First Paid Engagement

If you cannot speak fluently about all of this without notes, you're not ready for a paid engagement.

- **AI Governance fundamentals.** Define AI governance precisely and distinguish from data, IT, corporate governance. Explain all six core principles operationally. Resources: NIST AI RMF (nist.gov, 40 pages, free). EU AI Act executive summary.
- **Regulatory landscape.** Name Utah SB 149, Texas TRAIGA, Illinois HB 3773, Colorado SB 24-205, and the EU AI Act. Explain who each applies to and what it requires at a high level. Verify all regulatory details before client conversations.
- **Discovery session.** Run the full 60-minute session, score the maturity model live, produce the output template within 48 hours. Practice with two non-paying beta sessions before first paid engagement.
- **Objection handling.** Respond to all seven common objections without hesitation or script. Natural, not memorized. Say each one out loud five times.
- **Document anatomy.** Explain what every section of a complete AI policy covers and why it matters.
- **SMB positioning.** Articulate in 90 seconds what a diagnostic-led solo operator offers that larger firms cannot, using specific evidence.
- **SHIELD methodology.** Deliver the full methodology walkthrough in 5 minutes without reference.

8.2 Develop Through First 3-6 Months

These capabilities develop through real client exposure. Don't wait to accumulate before starting.

- **Industry-specific risk patterns.** After 5-10 clients, you will see patterns. Before that, research the industry before each engagement.
- **Vendor assessment depth.** Read three AI vendor contracts in full: OpenAI, Microsoft Copilot, Salesforce Einstein. Mark every clause on retention, training, incident notification.
- **Incident response design.** NIST SP 800-61 for the structural model. Adapt to AI context.
- **Policy language craft.** Write for specificity, not generality. Read three published AI policies from tech companies. Note what makes language specific vs generic.

- **Risk classification nuance.** Edge cases: tools that straddle tiers, tools low-risk individually but high-risk in aggregate, tools where risk depends on use rather than nature.
- **Content flywheel.** Produce consistent content while delivering. Build a capture system that logs insights from every client session for future content.

8.3 Long-Term Depth (12-24 Months)

These take time and become the moat.

- **ISO 42001 practitioner depth.** Purchase the ISO 42001 standard. Consider ISACA's implementation course if client demand warrants.
- **EU AI Act compliance advisory.** Study the Act article by article. Follow the European AI Office's guidance. Track Codes of Practice development.
- **Sector-specific expertise.** Develop deep competence in one or two sectors. Financial services: Federal Reserve SR 11-7. Healthcare: HHS AI guidelines. Professional services: ABA AI ethics opinions.
- **Algorithmic bias literacy.** "Fairness and Machine Learning" (Barocas, Hardt, Narayanan, fairmlbook.org, free). IBM AI Fairness 360 documentation.
- **Practitioner community standing.** IAPP AI governance community. ISACA AI Governance working group. Build referral flow, peer learning, early access to regulatory developments.

Toolkit Section 9: Failure Modes in Practice

These are the places the practice can fail. Name them. Design against them.

9.1 Governance Becomes a Dead-End Engagement

How it happens: Assessment is delivered. Client treats it as a completed project. No next conversation. Relationship dies.

Prevention:

- Build the Margin and Operational Risk Indicators section into every deliverable. Not optional.
- Schedule the Strategic Briefing at contract signature, not engagement end. Next conversation on the calendar before Assessment is delivered.
- Strategic Briefing has a structured format ending with "Here are the two ways we'd work together next." No ambiguity.
- Pre-write Sprint and Retainer engagement letters. Send within 48 hours of Strategic Briefing.

9.2 Credential Gap Triggers Buyer Objection

How it happens: Prospect Googles Boubacar, finds no IAPP certification, no ISO 42001, no law degree. Concludes this is amateur work.

Prevention:

- Manufacture authority proof that doesn't require credentials: case studies, named methodology (SHIELD), published content, speaking appearances, self-published book by month 12.
- Get ahead of the objection in the Qualification Call: "I'm not a lawyer. What I do is not a substitute for legal advice. What I do is give you the operational infrastructure that makes your lawyer's job easier and your vendor conversations sharper."
- Build a "Who We Work With" page on catalystworks.consulting with testimonials from beta and early clients.

9.3 Regulatory Hallucination

How it happens: Practitioner cites a law or regulation that doesn't apply or cites it incorrectly. Client acts on bad information. Liability.

Prevention:

- Verify every regulatory citation before a client conversation. Primary sources only (state legislative websites, EU official text).
- Every deliverable includes a clear disclaimer: "This output is a starting point for governance documentation, not legal advice. Any policy affecting employee rights, client data, or regulatory compliance should be reviewed by qualified legal counsel before implementation."
- Maintain a "last verified" date on any regulatory summary in templates and lead magnets.

9.4 Build-Instead-of-Outreach Pattern

How it happens: Practitioner builds infrastructure (agent, microsite, YouTube production workflow) as a substitute for the anxiety-inducing work of sending messages to specific humans. This has already happened (30 cold emails, zero replies, March 2026).

Prevention:

- Weekly outreach quota is non-negotiable. 20 governance emails, 5 LinkedIn DMs, 2 podcast pitches. Every week. Tracked in The Forge 2.0.

- Content shipping does not count as outreach. Building does not count as outreach. Only messages sent to specific humans asking for specific conversations count.
- Monthly review includes an honest audit: what percentage of my working hours went to outreach vs building? If under 20%, correct immediately.

9.5 Scope Creep Into Technical AI Design

How it happens: Client asks practitioner to advise on AI system architecture, model selection, bias mitigation. Practitioner takes the work. Work is outside competency. Bad outcomes.

Prevention:

- Scope the practice explicitly to governance documentation and advisory. Technical AI design is out of scope.
- When a question falls outside scope ("how do I make my model less biased?"), redirect: "That's a technical AI question and outside what I do. What I can do is name the governance requirements around bias monitoring and help you write the vendor questions to ask your AI provider about their bias testing."

9.6 Price Anchoring Failure

How it happens: Practitioner takes the first engagement at a discount to get it signed. Discount becomes the reference price. All future clients anchor to the discount. Practice becomes unprofitable.

Prevention:

- $3,500 is the Assessment price. Not negotiable.
- If the first prospect wants a discount, offer a beta engagement at no charge (no price anchor set) in exchange for a testimonial and a written case study.
- Never publish the Assessment price with a "from" or "starting at." Fixed price signals confidence. Variable pricing signals negotiation, which signals uncertainty.

Toolkit Section 10: The Governance Agent (Deferred)

The governance agent MVP is deferred to Phase 1 (post-first-revenue). Building a bot-driven diagnostic before the practice has paying clients is the build-instead-of-outreach failure mode in HD. When the agent is built, the following specifications apply.

10.1 What the Agent Must Do

- Collect organizational context (industry, employee count, jurisdictions, AI tool inventory, existing documentation, triggering question)
- Map jurisdiction to applicable regulations (Utah, Texas, Illinois, Colorado, EU)
- Walk organization through risk tier classification
- Generate governance policy sections reflecting specific context (not fill-in-the-blank templates)
- Run gap analysis against chosen framework (NIST RMF, ISO 42001, EU AI Act)
- Explain governance concepts in plain language

10.2 What the Agent Must NOT Do

- Generate full compliance opinions for specific regulations
- Provide legal advice or claim to substitute for legal review
- Generate incident response plans (too context-specific for v1)
- Perform bias audits or technical AI assessments
- Advise on AI system design, model selection, technical architecture

10.3 Design Safeguards

- Every generated document includes an explicit completeness statement listing what is and isn't covered
- Every regulatory claim has a verifiable citation
- Every output includes the non-buried disclaimer about legal advice
- Agent is scoped to governance only; technical questions redirect
- Knowledge freshness flag on all regulatory content with "last verified" date
- Human review gate on any regulatory content update

10.4 V1 Expansion Criteria

V2 unlocks only after: at least 50 organizations have completed v1 intake, output quality validated against real client situations, legal review has confirmed scope and disclaimer adequacy, regulatory update mechanism is operational.

Toolkit Appendix A: Regulatory Quick Reference (Verified April 2026)

Jurisdiction	Law	Effective Date	Primary Obligation	Penalty
Utah	SB 149 AI Policy Act	May 1, 2024	Disclose generative AI to consumers on request; responsible for AI-driven deception	Up to $2,500 per violation
Texas	HB 149 TRAIGA	Jan 1, 2026	Intent-based prohibition on AI misuse (discrimination, manipulation, deepfakes, exploitation)	Up to $200,000 per violation; 60-day cure period
Illinois	HB 3773	Jan 1, 2026	Prohibit AI-driven employment discrimination; disclose AI use in employment decisions	Civil rights violation remedies
Colorado	SB 24-205 AI Act	Jun 30, 2026	Reasonable care against algorithmic discrimination for high-risk AI deployers and developers	Subject to AG rulemaking; under political revision
EU	AI Act	In force Aug 1, 2024; high-risk Aug 2, 2026 (may be delayed)	Risk-tiered obligations; prohibited practices banned; high-risk systems require conformity assessment	Up to €35M or 7% turnover

Always verify against primary sources before client conversations. This table is a reference, not a substitute for current verification.

Toolkit Appendix B: 7-Day Actions at Each Stage

Stage 1 (Weeks 1-2): Write and publish first LinkedIn post this week. Topic: "What AI governance is, and what it is not." Target: 400-600 words.

Stage 2 (Weeks 2-3): Book two beta diagnostic sessions with people in the network. Frame: "I'm doing research on how organizations are managing AI right now. I'd like to run a 45-minute diagnostic and give you a written summary. No charge, no pitch." Run both.

Stage 3 (Weeks 3-4): Write out personal answers to "We're too small to need this" and "We'll deal with it when regulations require it." Say them aloud until natural. Use in next three prospecting conversations.

Stage 4 (Weeks 4-5): Draft the three-tier roadmap as a one-page visual. Last slide of every discovery session: "Here is where you are and where this goes."

Stage 5 (Weeks 5-8): Write a one-page scope document for the governance agent (deferred to Phase 1): what it does, what it does not do, three failure modes designed against. Clarity document, not technical spec.

End of playbook.